

# COMP3411/9814: Artificial Intelligence

## Planning

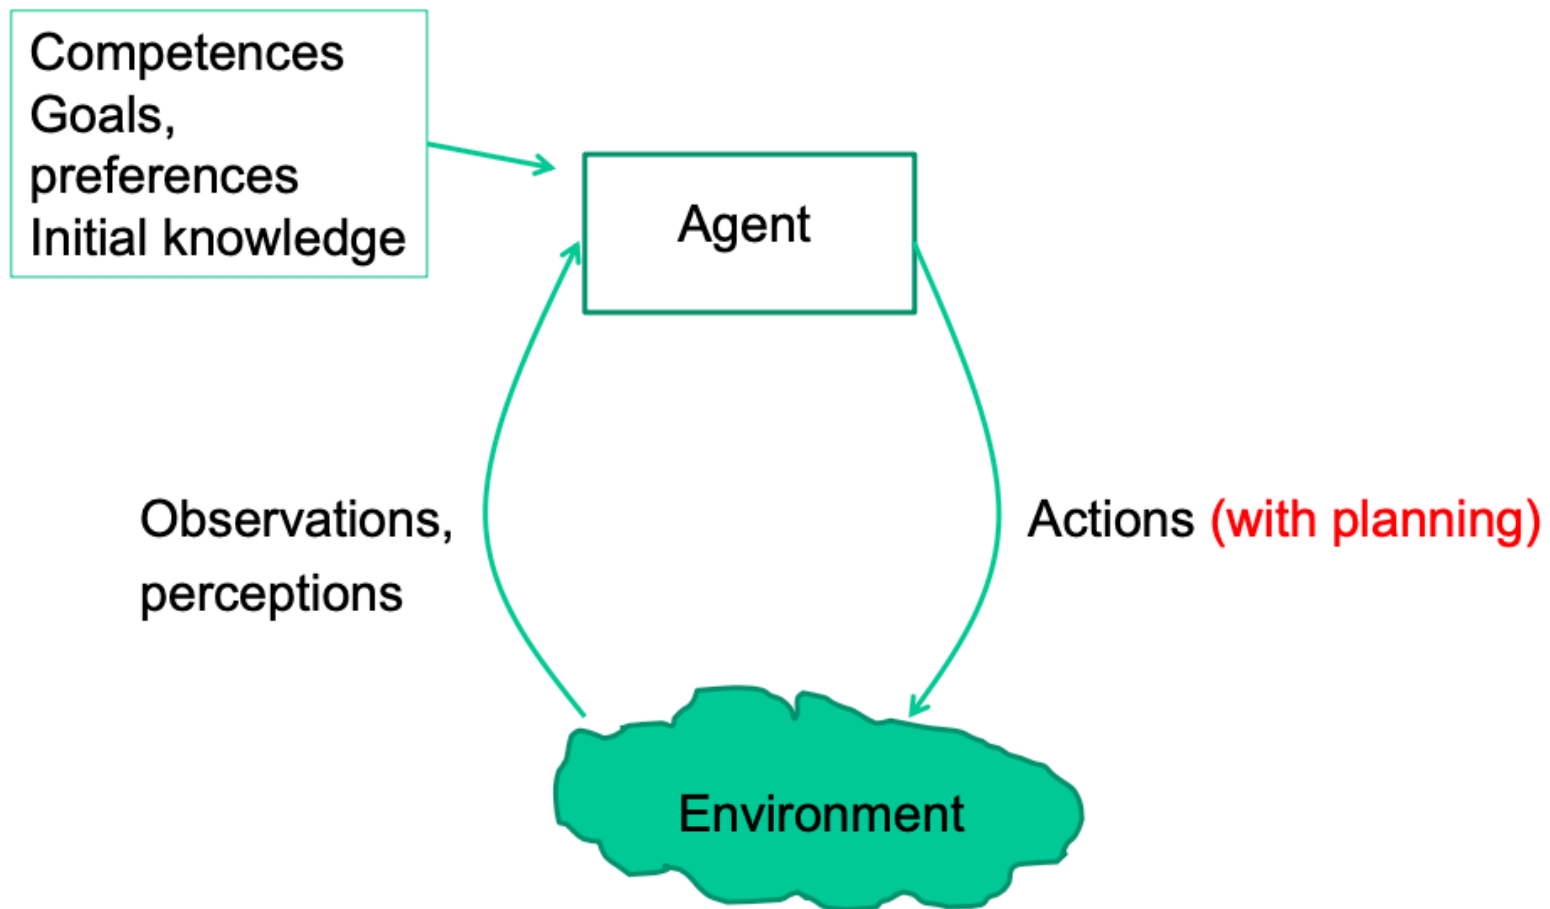
# Lecture Overview

---

- ❑ Reasoning About Action
- ❑ STRIPS Planner
- ❑ Forward planning
- ❑ Regression Planning
- ❑ GraphPlan
- ❑ Planning as Constraint Satisfaction

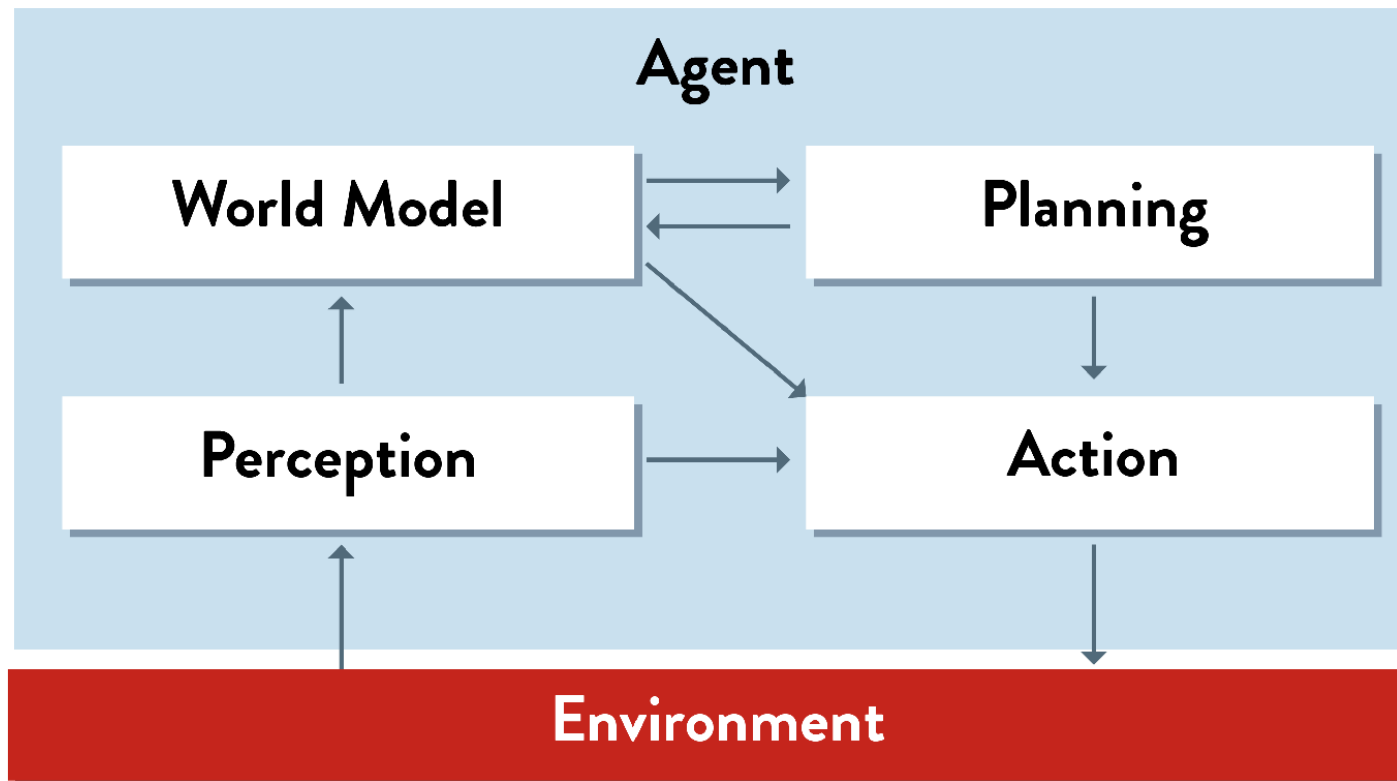
# Agent acting in its environment

---



# Planning Agent

---



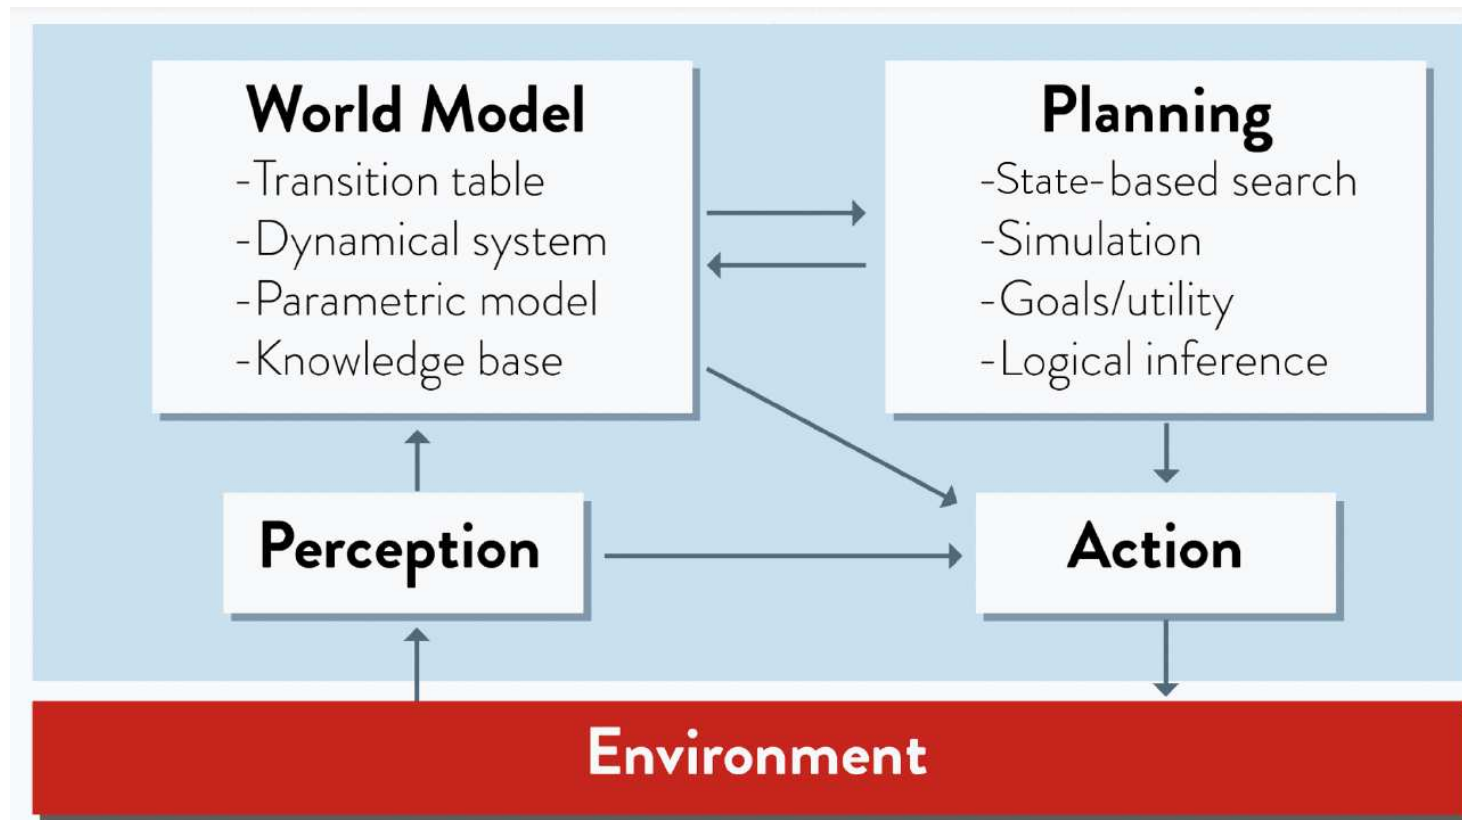
Goal-Based Agent

# Planning Agent

---

- ❑ Decision making of this kind is fundamentally different from the condition– action rules
- ❑ It involves consideration of the future
  - “What will happen if I do such-and-such?” and
  - “Will that make me happy?”
- In the reflex agent designs, this information is not explicitly represented

# Models and Planning



# Agent Plans Actions To Achieve Desired Goals

---

- ❑ “PLANNING” in general sense includes problem solving in state space
- ❑ “PLANNING” in narrow sense is: “means ends planning”

# Planning

---

- ❑ Planning is deciding what to do based on an agent's ability, its goals. and the state of the world.
- ❑ Planning is finding a sequence of actions to solve a goal. Initial assumptions:
  - The world is deterministic.
  - There are no exogenous events outside of the control of the robot that change the state of the world.
  - The agent knows what state it is in.
  - Time progresses discretely from one state to the next.
  - Goals are predicates of states that need to be achieved or maintained.



# Planning Agent

---

- ❑ The planning agent or goal-based agent is more flexible because the knowledge that supports its decisions is represented explicitly and **can be modified**.
- ❑ The agent's behavior can easily be changed.

# Planning Agent

---

- ❑ Environment changes due to the performance of actions
- ❑ Planning scenario
  - Agent can control its environment
  - Only atomic actions, not processes with duration
  - Only single agent in the environment (no interference)
  - Only changes due to agent executing actions (no evolution)
- ❑ More complex examples
  - Robocup dog
  - Delivery robot
  - Self-driving car

# Representation

---

□ How to represent a classical planning problem?

# Representation

---

- ❑ How to represent a classical planning problem?
- ❑ Representing with States, Actions, and Goals


# Actions

---

- ❑ A deterministic **action** is a partial function from states to states.
- ❑ The **preconditions** of an action specify when the action can be carried out.
- ❑ The **effect** of an action specifies the resulting state.

# Delivery Robot Example

---




The delivery robot domain

The robot, called Rob, can buy coffee at the coffee shop, pick up mail in the mail room, move, and deliver coffee and/or mail.

# Delivery Robot Example

---

**Features:**


*RLoc* – Rob's location  
*RHC* – Rob has coffee  
*SWC* – Sam wants coffee  
*MW* – Mail is waiting  
*RHM* – Rob has mail

**Actions:**

*mc* – move clockwise  
*mcc* – move counterclockwise  
*puc* – pickup coffee  
*dc* – deliver coffee  
*pum* – pickup mail  
*dm* – deliver mail

# Delivery Robot Example

---

**Features:**

*RLoc* – Rob's location  
*RHC* – Rob has coffee  
*SWC* – Sam wants coffee  
*MW* – Mail is waiting  
*RHM* – Rob has mail

Features to describe states

**Actions:**

*mc* – move clockwise  
*mcc* – move counterclockwise  
*puc* – pickup coffee  
*dc* – deliver coffee  
*pum* – pickup mail  
*dm* – deliver mail

Robot actions



## *State description*

---

*The state is described in terms of the following features:*

- RLoc - the robot's location, which is one of the coffee shop (*cs*), Sam's office (*off*), the mail room (*mr*) or in or the laboratory (*lab*)
- SWC - Sam wants coffee. The atom *swc* means Sam wants coffee and  $\neg \text{swc}$  means Sam does not want coffee.

# Robot Actions

---

## □ Rob has six actions

- Rob can move clockwise (*mc*)
- Rob can move counterclockwise (*mcc*) or (*mac*), for now we use (*mcc*).
- Rob can pick up coffee if Rob is at the coffee shop.
  - *puc* mean that Rob picks up coffee.
- .....

## □ Assume that it is only possible for Rob to do one action at a time.

# Explicit State-space Representation

---

# Explicit State-space Representation

---

- ❑ The states are specifying the following:
  - the robot's location,
  - whether the robot has coffee,
  - whether Sam wants coffee,
  - whether mail is waiting,
  - whether the robot is carrying the mail.

$\langle lab, \neg rhc, swc, \neg mw, rhm \rangle$

# Explicit State-space Representation

<i>State</i>	<i>Action</i>	<i>Resulting State</i>
$\langle lab, \neg rhc, swc, \neg mw, rhm \rangle$	<i>mc</i>	$\langle mr, \neg rhc, swc, \neg mw, rhm \rangle$
$\langle lab, \neg rhc, swc, \neg mw, rhm \rangle$	<i>mcc</i>	$\langle off, \neg rhc, swc, \neg mw, rhm \rangle$
$\langle off, \neg rhc, swc, \neg mw, rhm \rangle$	<i>dm</i>	$\langle off, \neg rhc, swc, \neg mw, \neg rhm \rangle$
$\langle off, \neg rhc, swc, \neg mw, rhm \rangle$	<i>mcc</i>	$\langle cs, \neg rhc, swc, \neg mw, rhm \rangle$
$\langle off, \neg rhc, swc, \neg mw, rhm \rangle$	<i>mc</i>	$\langle lab, \neg rhc, swc, \neg mw, rhm \rangle$
...	...	...

The complete representation includes the transitions for the other 62 states.

# Explicit State-space Representation

---

This is not a good representation:

- ❑ There are usually too many states to represent, to acquire, and to reason with.
- ❑ Small changes to the model mean a large change to the representation.
  - Adding another feature means changing the whole representation.
- ❑ It does not represent the structure of states;
  - there is much structure and regularity in the effects of actions that is not reflected in the state transitions.

# STRIPS language for problem definition

---

- ❑ STRIPS=Stanford Research Institute Problem Solver
- ❑ STRIPS—traditional representation  
“STRIPS-like representation”
- ❑ STRIPS makes some simplifications:
  - no variables in goals
  - positive relations given only
  - unmentioned relations are assumed false (c.w.a. – closed world assumption)
  - effects are conjunctions of relations

# STRIPS Representation

---

- ❑ Divide the features into:

- primitive features
- derived features. There are rules specifying how derived can be derived from primitive features.

- ❑ For each action:

- **precondition** that specifies when the action can be carried out.
- **effect** a set of assignments of values to primitive features that are made true by this action.

STRIPS assumption: every primitive feature not mentioned in the effects is unaffected by the action.



# Example STRIPS representation

---

Pick-up coffee (puc):

- **precondition**: [cs,¬rhc]
- **effect**: [rhc]

Deliver coffee (dc):

- **precondition**: [off,rhc]
- **effect**: [¬rhc,¬swc]

# Feature-based representation of actions

---

- ❑ STRIPS is an action-centric representation
- ❑ A feature-centric representation is more flexible, as it allows for conditional effects, and non-local effects.

# Feature-based representation of actions

---

- For each action:

- **precondition** is a proposition that specifies when the action can be carried out.

- For each feature:

- **causal rules** that specify when the feature gets a new value and
- **frame rules** that specify when the feature keeps its value.

# Example feature-based representation

---

- ❑ Precondition of pick-up coffee (puc):

$$RLoc=cs \wedge \neg rhc$$

- ❑ Rules for location is cs:

$$RLoc'=cs \leftarrow Rloc = off \wedge Act=mcc$$

$$RLoc'=cs \leftarrow Rloc = mr \wedge Act=mc$$

$$RLoc'=cs \leftarrow Rloc = cs \wedge Act \neq mcc \wedge Act \neq mc$$

- ❑ Rules for “robot has coffee”

$$rhc' \leftarrow rhc \wedge Act \neq dc$$

$$rhc' \leftarrow Act=puc$$

# Example feature-based representation

---

- Precondition of pick-up coffee (puc):

$$RLoc=cs \wedge \neg rhc$$

- Rules for location is cs:

$$RLoc'=cs \leftarrow Rloc = off \wedge Act=mcc$$

$$RLoc'=cs \leftarrow Rloc = mr \wedge Act=mc$$

$$RLoc'=cs \leftarrow Rloc = cs \wedge Act \neq mcc \wedge Act \neq mc$$

causal rules

frame rule

- Rules for “robot has coffee”


$$rhc' \leftarrow rhc \wedge Act \neq dc$$

$$rhc' \leftarrow Act=puc$$

# States are Represented with Relations

---

Example state from blocks world




This state can be represented by the following set of relations:

# States are Represented with Relations

---

Example state from blocks world



This state can be represented by the following set of relations:

**on(c,a), on(a,1), on(b,3), clear(2), clear(4), clear(b), clear(c)**

# Defining Goals and possible Actions

---

- Example of goals:


**$\text{on}(a,b)$ ,  $\text{on}(b,c)$**

- Example of action:

**$\text{move}(a, 1, b)$**

(Move block a from 1 to b)

- Action preconditions:





# Defining Goals and possible Actions


- Example of goals:

**$\text{on}(a,b), \text{on}(b,c)$**

- Example of action:

**$\text{move}(a, 1, b)$**

(Move block a from 1 to b)



- Action preconditions:

**$\text{clear}(a), \text{on}(a,1), \text{clear}(b)$**

- Action effects:

**$\text{on}(a,b), \text{clear}(1), \sim\text{on}(a,1), \sim\text{clear}(b)$**

“add” (true after action)

“delete” (no longer true after action)

( $\sim = \neg$ )

# Action Schema

---

- ❑ **Action schema** represents a set of actions using variables (variable names here written with capital initials)

**move( X, Y, Z)**

X is any block

Y and Z are any block or location

- ❑ **Precondition:** `on(X,Y),clear(X),clear(Z)`
- ❑ **Adds:** `on(X,Z),clear(Y)`
- ❑ **Deletes:** `on(X,Y),clear(Z)`

# STRIPS language for problem definition

---

□ STRIPS makes some simplifications:

- no variables in goals
- positive relations given only
- unmentioned relations are assumed false (c.w.a. – closed world assumption)
- effects are conjunctions of relations

# ADL - Action Description Language

- ❑ ADL removes some of the STRIPS assumptions, for example:

STRIPS	ADL
<i>States: + literals only</i> <b>on(a,b), clear(a)</b>	<b>on(a,b), clear(a), ~clear(b)</b>
<i>Effects: + literals only</i> Add <b>clear(b)</b> , Delete <b>clear(c)</b>	Add <b>clear(b)</b> and <b>~clear(c)</b> Delete <b>~clear(b)</b> and <b>clear(c)</b>
<i>Goals: no variables</i> <b>on(a,c), clear(a)</b>	<b>Exists X: on(X,c), clear(X)</b>

# STRIPS Representation

---

- ❑ Divide the features into:

- primitive features
- derived features. There are rules specifying how derived can be derived from primitive features.

- ❑ For each action:

- **precondition** that specifies when the action can be carried out.
- **effect** a set of assignments of values to primitive features that are made true by this action.


STRIPS assumption: every primitive feature not mentioned in the effects is unaffected by the action.

# Domain Specification For Blocks World

---

Action:

**move( Block, From, To)**



Action precondition:

**clear( Block), clear( To), on( Block, From)**

Positive effects (“add”):

**on(Block,To), clear(From)**

Negative effects (“del”)


**on(Block,From), clear(To)**

# Better With Additional Constraints

---

Action:

**move( Block, From, To)**



Precondition for Action:

**clear( Block), clear( To), on( Block, From)**

Additional constraints:

block( Block),	% Object Block to be moved must be a block
object( To),	% "To" is an object, i.e. a block or a place
To $\neq$ Block,	% Block cannot be moved to itself
object( From),	% "From" is a block or a place
From $\neq$ To,	% Move to new position
Block $\neq$ From	

# Planning

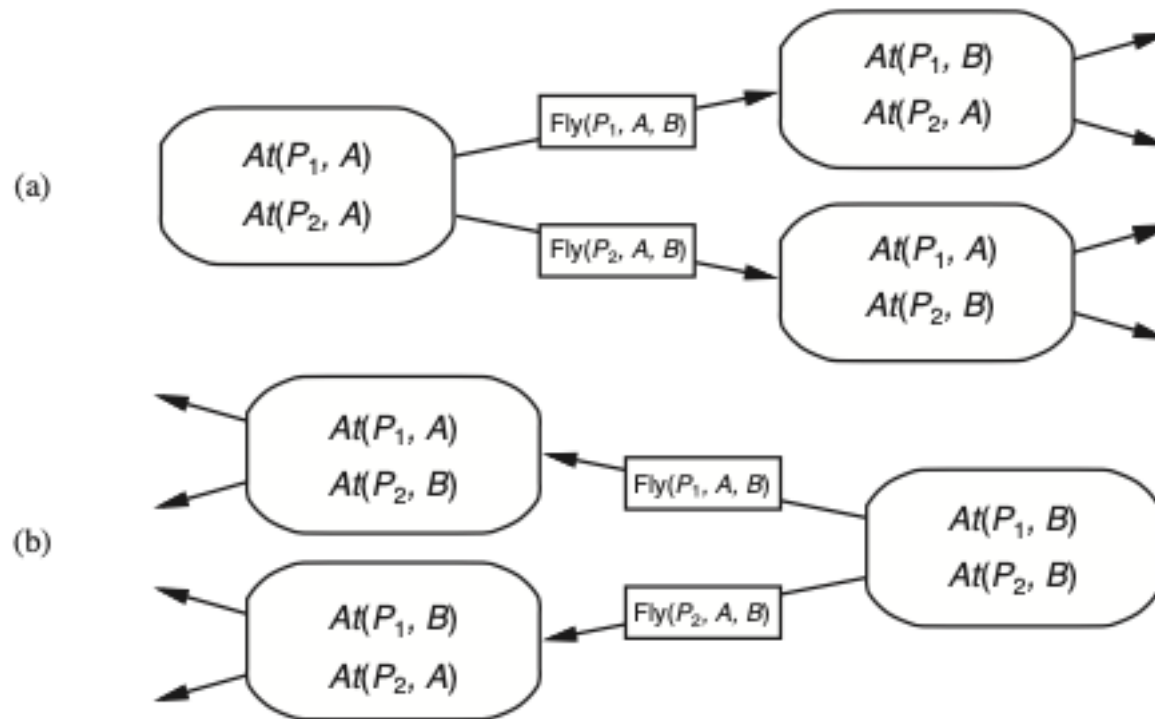
---

- ❑ Plan — sequence (or ordered set) of actions to achieve some goal
- ❑ Planner — problem solver that produces plans
- ❑ Goal — typically a conjunction of literals
- ❑ Initial State — typically a conjunction of literals
- ❑ Blocks World Example for goal  $on(B, C) \wedge on(C, Table)$ 
  - $move(C, A, Table), move(B, Table, C)$



# Simple Planning Algorithms

## □ Forward search and goal regression



Problem with forward search is state space can be very large

Problem with regression is that it is hard and doesn't always work

# Forward Search with Plan Graphs

---

- ❑ Only consider “propositional” plans
  - $S_i$  contains all literals that *could* hold at time  $i$
  - $A_i$  contains all actions that *could* have preconditions satisfied at time  $i$
  - Actions linked to preconditions
  - Literals that persist from time  $i$  to time  $i + 1$  linked via actions
  - Mutual exclusion (mutex) links between actions/literals at same time

# Mutual Exclusion

---

## □ Actions

- **Inconsistent effects:** One action negates an effect of the other
- **Competing needs:** Precondition of one action is mutually exclusive with a precondition of the other

## □ Literals

- One literal is the negation of the other
- **Inconsistent support:** Each possible pair of actions that could achieve the two literals is mutually exclusive

# Specification Of Blocks And Locations

---

% Our blocks world: three blocks a, b and c, and 4 locations  
block( a). block( b). block( c).

place( 1). place( 2). place( 3). place( 4).

% X is an object if X is a block or a place

object( X)  $\leftarrow$  ( block(X)  $\vee$  place(X) )

# Robots On Grid In Strips

---

4	5	6
<b>a</b> 1	<b>b</b> 2	<b>c</b> 3

Robots: a, b, c, cells 1, ..., 6

Goal: at(a,3)

Plan:  $m(b,2,5) \longrightarrow m(a,1,2) \longrightarrow m(c,3,6) \longrightarrow m(a,2,3)$

# GraphPlan Algorithm

---

Graph = Initial plan graph with initial state  $S_0$

nogoods = empty set

For  $t = 0, \dots$

- ▶ If all goals are non-mutex in  $S_t$ 
  - Extract solution from graph
    - Graph as CSP with variables T/F for when action in the plan
    - Or heuristically guided regression from  $S_t$  to  $S_0$
  - If valid solution, return solution
- ▶ If graph and nogoods didn't change then return failure
- ▶ Expand graph to next level

# GraphPlan Expansion Step

---

Add actions to  $A_i$  whose preconditions are at  $S_i$

Add “persistence actions” to  $A_i$  for literals from  $S_i$

Add mutex links to  $A_i$  for actions that cannot occur together

Add effects of all actions in  $A_i$  to  $S_{i+1}$

Add literals to  $S_{i+1}$  for persistence actions from  $A_i$

Add mutex links to  $S_{i+1}$  for literals that cannot occur together

# Forward Search with Plan Graphs

---

- ❑ A forward planner searches the state-space graph from the initial state looking for a state that satisfies a goal description.
- It can use any of the search strategies



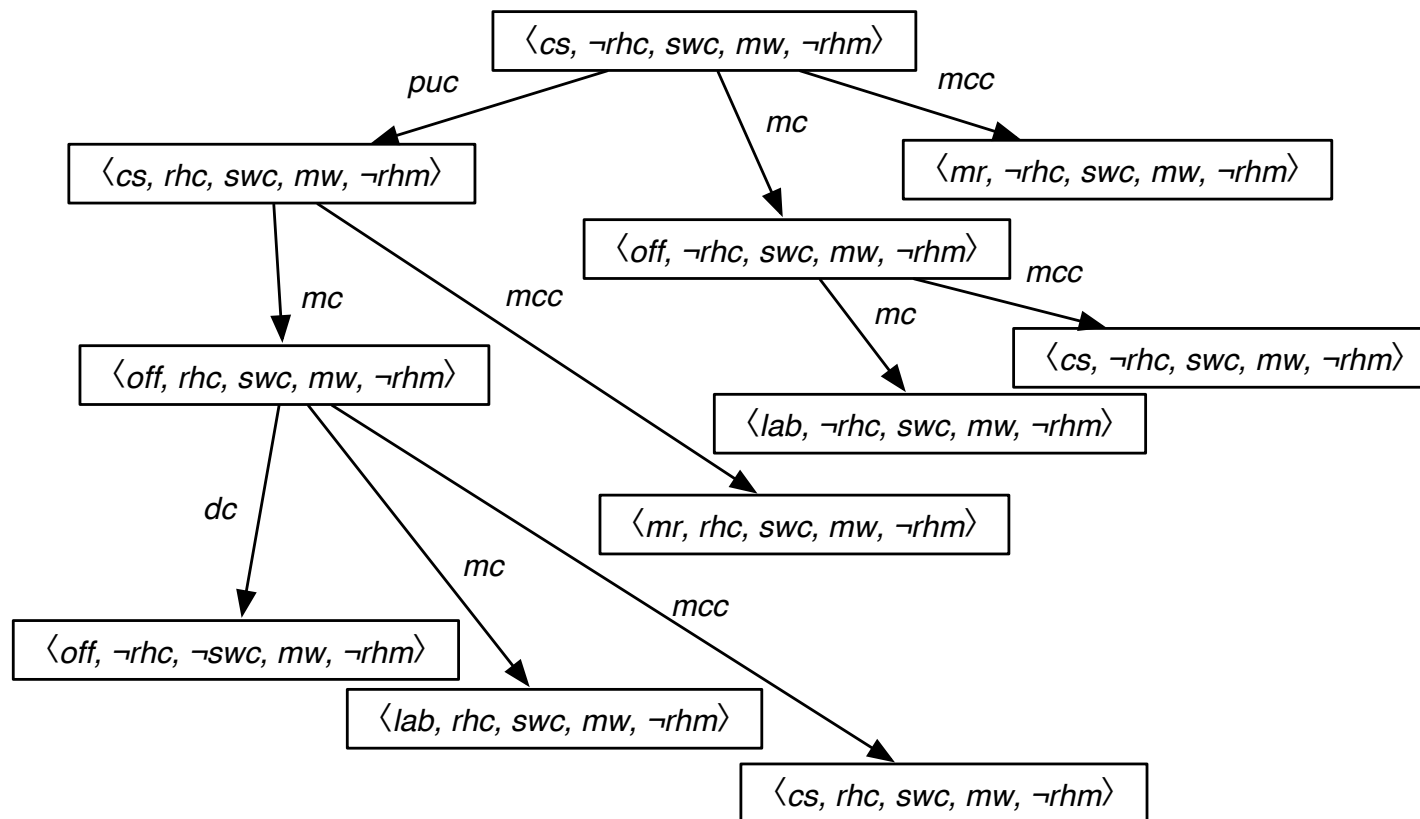
# Forward Search with Plan Graphs

---

The search graph is defined as follows:

- ❑ The nodes are states of the world, where a state is a total assignment of a value to each feature.
- ❑ The arcs correspond to actions.
- ❑ The start node is the initial state.
- ❑ The goal condition for the search, *goal(s)* is true if state *ss* satisfies the achievement goal.
- ❑ A path corresponds to a plan that achieves the goal.

# Forward Search with Plan Graphs



# Forward Search with Plan Graphs

## Actions

*mc*: move clockwise

*mac*: move anticlockwise

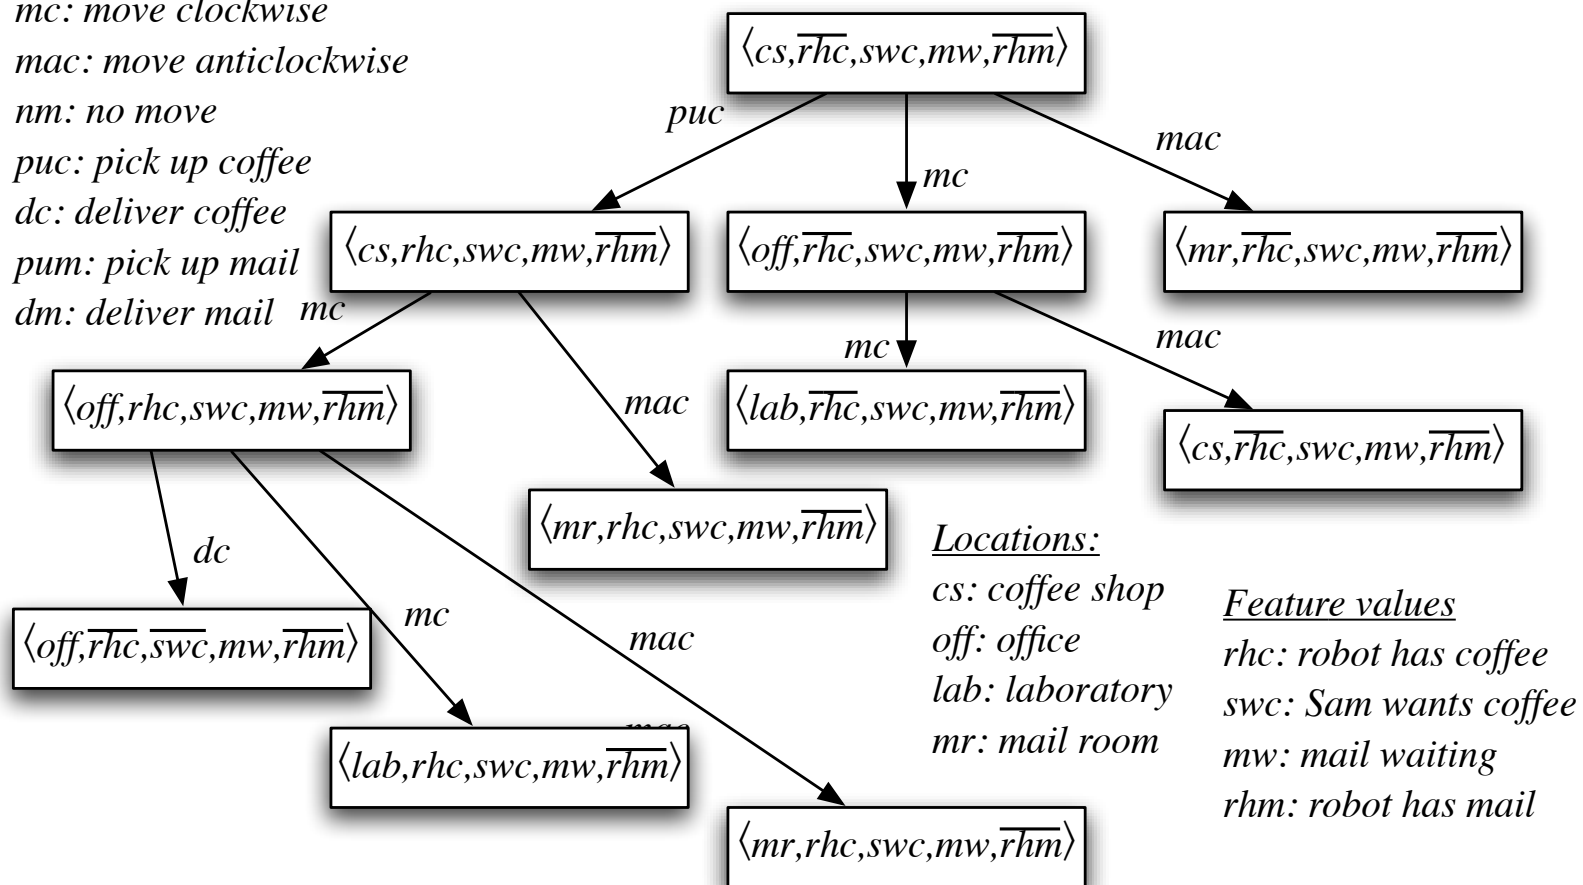
*nm*: no move

*puc*: pick up coffee

*dc*: deliver coffee

*pum*: pick up mail

*dm*: deliver mail




# Forward Search with Plan Graphs

---

- ❑ Using a forward planner is not the same as making an explicit state-based representation of the actions
- ❑ The relevant part of the graph is created dynamically from the representations of the actions.

## Principle of means-ends analysis (Optional)

---



In this state, the following relations hold:

**on(c,a), on(a,1), on(b,3), clear(2), clear(4), clear(b), clear(c)**

Let goal of plan be **on(a,b)**; find plan:

What action establishes **on(a,b)**? Such action is: **move(a,X,b)**


What is precondition COND for this action?

COND: **on(a,X), clear(a), clear(b)**

Set intermediate goal COND, find plan for COND Etc.

## Principle of means-ends analysis (Optional)

---



In this state, the following relations hold:

**on(c,a), on(a,1), on(b,3), clear(2), clear(4), clear(b), clear(c)**


Let goal of plan be **on(a,b)**; find plan:

What action establishes **on(a,b)**? Such action is: **move(a,X,b)**

What is precondition COND for this action?

## Principle of means-ends analysis (Optional)

---



In this state, the following relations hold:

**on(c,a), on(a,1), on(b,3), clear(2), clear(4), clear(b), clear(c)**

Let goal of plan be **on(a,b)**; find plan:

What action establishes **on(a,b)**? Such action is: **move(a,X,b)**


What is precondition COND for this action?

COND: **on(a,X), clear(a), clear(b)**

Set intermediate goal COND, find plan for COND Etc.

## Principle of means-ends analysis (Optional)

---



In this state, the following relations hold:

**on(c,a), on(a,1), on(b,3), clear(2), clear(4), clear(b), clear(c)**

Let goal of plan be **on(a,b)**; find plan:

What action establishes **on(a,b)**? Such action is: **move(a,X,b)**

What is precondition COND for this action?

COND: **on(a,X), clear(a), clear(b)**

Set intermediate goal COND, find plan for COND Etc.



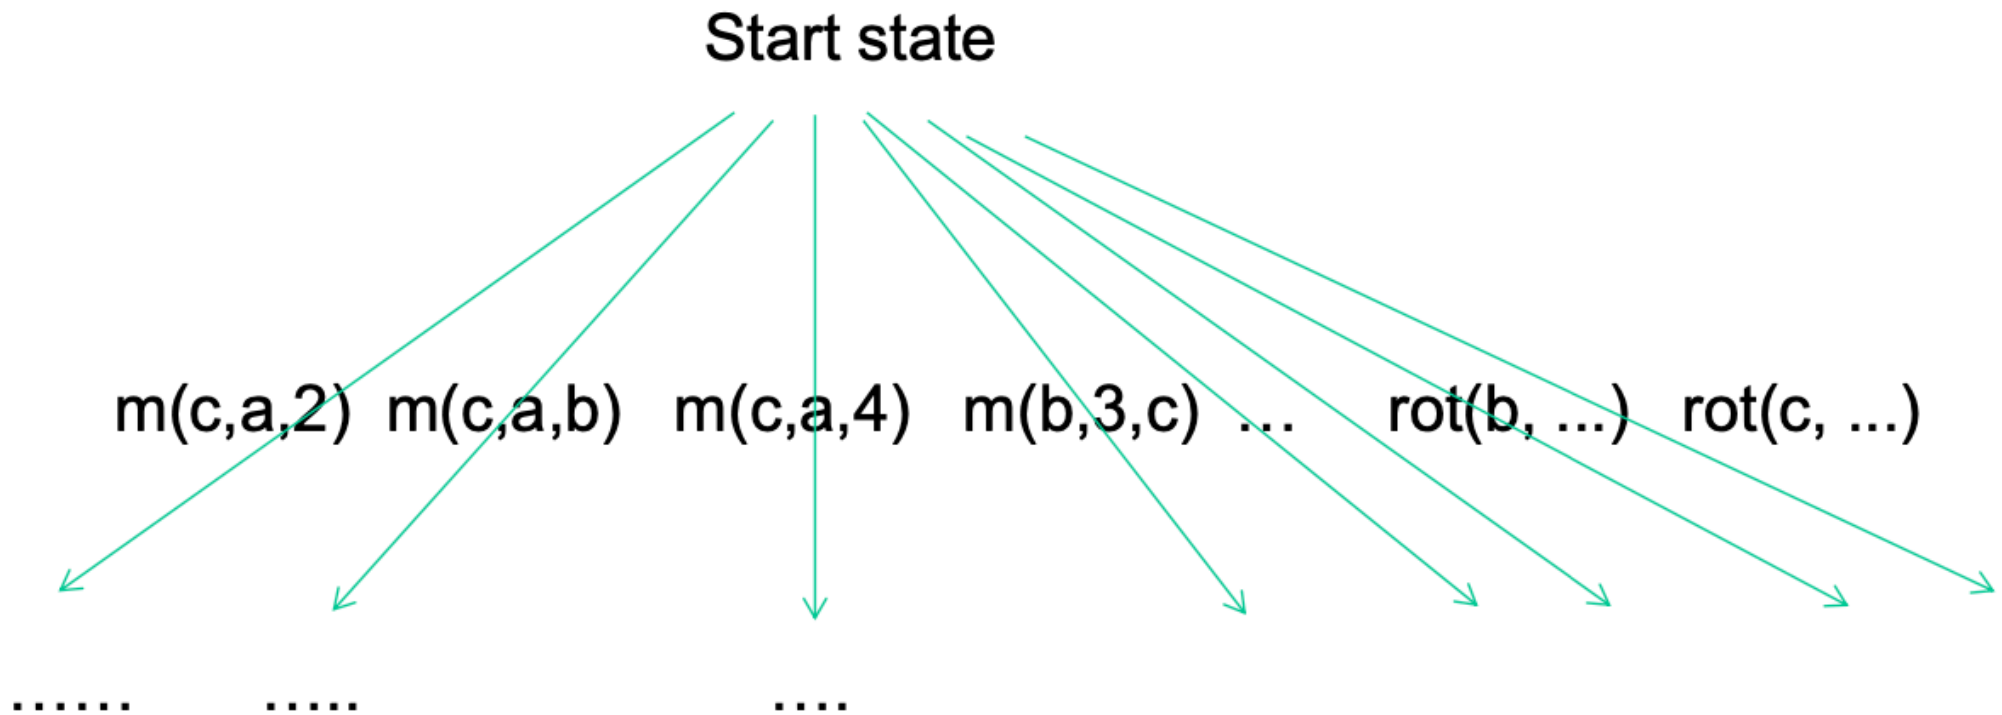
# Comparison with state space (Optional)

---

- ❑ Instate-space: search state space
- ❑ In means-ends planning: search space of sets of goals
- ❑ Space of sets of goals = abstraction of state space
- ❑ What is better? Means-ends planning may be able to avoid searching useless actions, see example on next slide

# Actions in state space (Optional)


---



# Means-ends planning in Strips (Optional)

---

One possible realisation of means-ends planning



# Nondeterministic strips algorithm

---

```
procedure plan( InitialState, Goals, Plan, FinalState)
  if Goals  $\subseteq$  InitialState then Plan = [ ] else           % All goals achieved
  begin
    Select a goal G from Goals;
    Select an action A that achieves G;                     % adds( A, G)
    PreCond = preconditions of A;
    plan( InitialState, PreCond, PrePlan, MidState1) ;      % Enable A
    Apply A to MidState1 giving MidState2;
    plan( MidState2, Goals, PostPlan, FinalState);          % Achieve remaining goals
    Plan = concatenate( PrePlan, [ Action ], PostPlan)
  end
```

# Strips in Prolog

---

**% plan( State, Goals, Plan, FinalState)**

**plan( State, Goals, [ ], State) :-  
    satisfied( State, Goals).**

**plan( State, Goals, Plan, FinalState) :-**

**conc( PrePlan, [Action | PostPlan], Plan),**

**select( State, Goals, Goal),**

**achieves( Action, Goal)**

**can( Action, Condition),**

**plan( State, Condition, PrePlan, MidState1),**

**apply( MidState1, Action, MidState2),**

**plan( MidState2, Goals, PostPlan, FinalState).**

% Divide plan

% Select a goal

% Relevant action

% Enable Action

% Apply Action

% Remaining goals

# Additional details

---

- ❑ Search strategy (depth-first, breadth-first,...)
- ❑ Goal protection: do not destroy what you already achieved!
- ❑ But: goal protection is not always possible!

# Realisation with iterative deepening

---

- ❑ Start with plan of length 0 and keep increasing maximal allowed length of plan, until plan is found
- ❑ On each iteration (for each maximal plan length) search all possible plans with depth-first search

# Realisation with iterative deepening


---

- ❑ Start with plan of length 0 and keep increasing maximal allowed length of plan, until plan is found
- ❑ On each iteration (for each maximal plan length) search all possible plans with depth-first search
- ❑ Surprise is possible, for example in the case of Sussman's anomaly (on next slide)



# Sussman's anomaly

---



Goals:  $\text{on}(a,b)$ ,  $\text{on}(b,c)$

Basic STRIPS planner with breadth-first search produces:

move( c, a, 2)

move( b, 3, a)

move( b, a, c)

move( a, 1, b)

??? What is the point of this ???

Problem is: STRIPS concentrates on solving a single goal at a time

# Completeness


---

- ❑ Even with global iterative deepening, our planner still has problems.
- ❑ E.g. it finds a four step plan above for our example block stack
- ❑ Why STRIPS cannot find the optimal, 3-step plan? Basic STRIPS is **incomplete**! It does not consider all possible plans.
- ❑ Problem: locality (only work towards achieving one goal G at a time, temporarily ignoring other goals until G is achieved)
- ❑ Sometimes referred to as “linearity” (goals are achieved in “linear order”)

# Goal regression

---

- ❑ STRIPS solves goals one after another “locally” (when solving one goal it does not consider other goals)
- ❑ Better: “global planning” (keep in mind all the goals all the time)
- ❑ One idea to achieve global planning is **goal regression**
- ❑ This is based on concept of “Regressing Goals through Action”



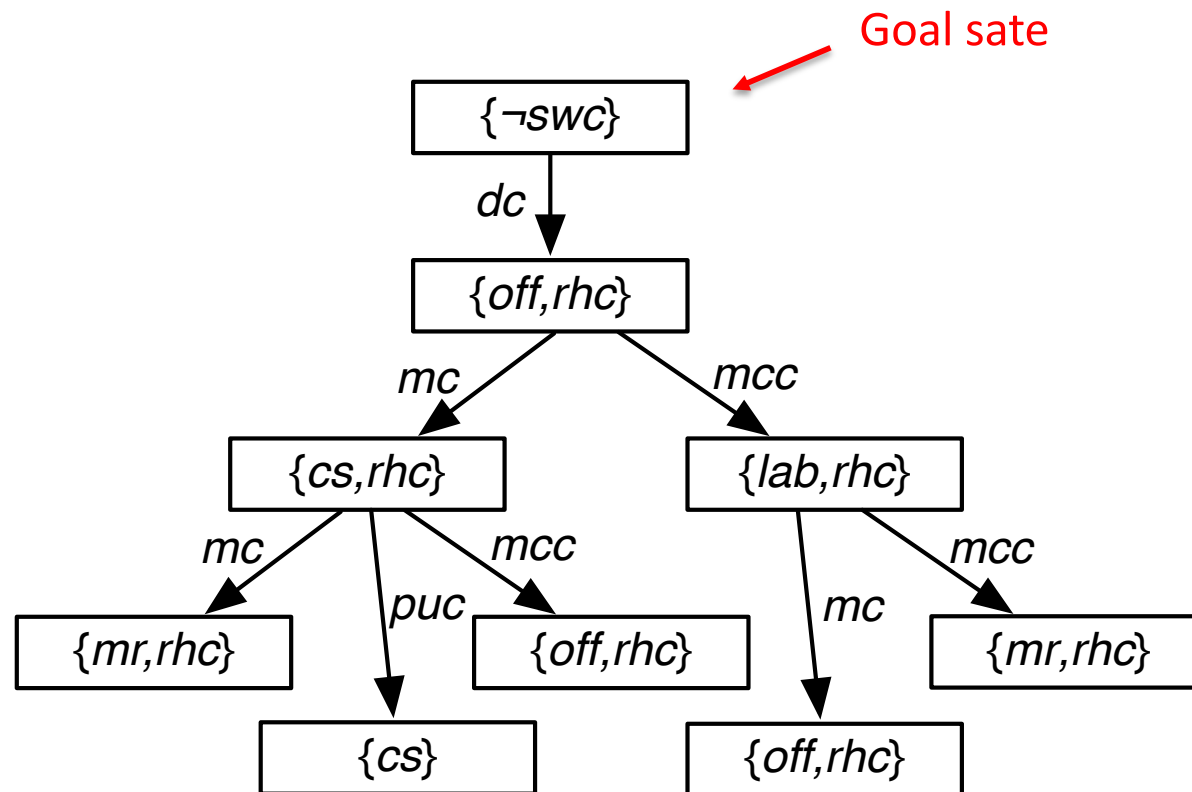
# Goal regression

---

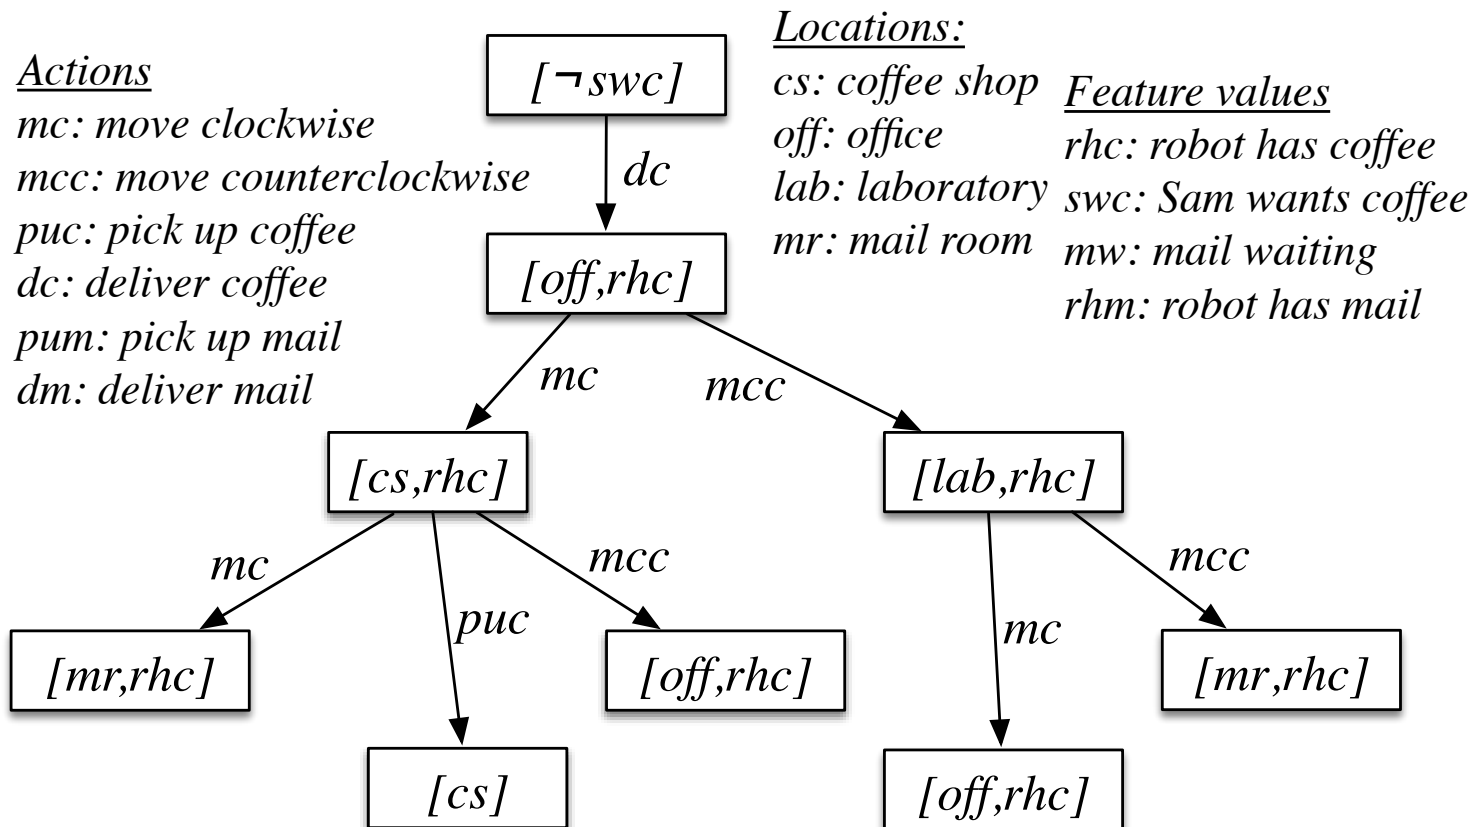
**Regression planning** is searching in the graph defined by the following:

- ❑ The nodes are subgoals.
- ❑ The arcs correspond to actions. An arc from node  $g$  to  $g'$ , labeled with action  $act$ , means
  - $act$  is the last action that is carried out before subgoal  $g$  is achieved, and
  - node  $g'$  is a subgoal that must be true immediately before  $act$  so that  $g$  is true immediately after  $act$ .
- ❑ The start node is the planning goal to be achieved.
- ❑ The goal condition for the search,  $goal(g)$ , is true if  $g$  is true of the initial state.

# Goal regression




# Goal regression



# Goal regression

---




$$\text{RegressedGoals} = \text{Goals} + \text{can}(A) - \text{add}(A)$$

Goals and  $\text{del}(A)$  are disjoint

# Goal regression

---



Goal regression enables “global” planning:

Planner can see all relevant goals at any point of planning



# Robots On Grid In Strips

---

4	5	6
<b>a</b> 1	<b>b</b> 2	<b>c</b> 3

Robots: a, b, c,      Cells 1, ..., 6

Goal: at(a,3)

# Domain definition

---

Action: Robot R moves from A to B

$m(R,A,B)$

Preconditions:

$at(R,A), c(B)$

Additional constraints:

$robot(R), adjacent(A,B)$       % R is a robot, A and B are adjacent

Positive effects:

$at(R,B), c(A)$

Negative effects:

$at(R,A), c(B)$

# Robots On Grid In Strips

---

4	5	6
<b>a</b> 1	<b>b</b> 2	<b>c</b> 3

Robots: a, b, c, cells 1, ..., 6

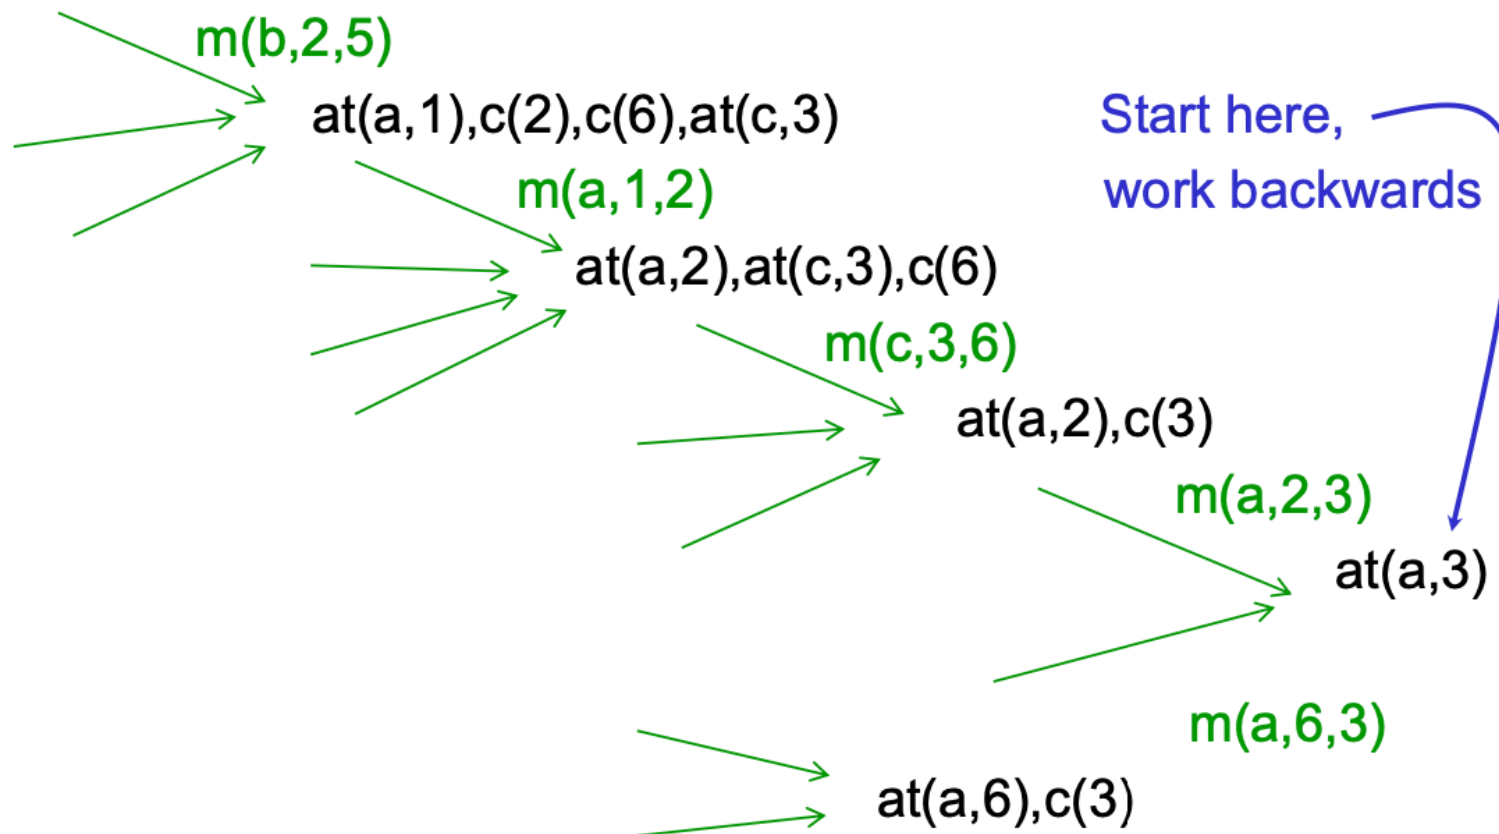
Goal: at(a,3)

Plan:  $m(b,2,5) \longrightarrow m(a,1,2) \longrightarrow m(c,3,6) \longrightarrow m(a,2,3)$

# Finding a plan for goal $\text{at}(a,3)$


Initial state:  $\text{at}(a,1), \text{at}(b,2), \text{at}(c,3), c(4), c(5), c(6)$

$\text{at}(b,2), c(5), \text{at}(a,1), c(6), \text{at}(c,3)$



# State space for planning with goal regression

---



# Goal state and heuristic

---

- ❑ “Goal” condition for this search space:  
RegressedGoals is a subset of StartState
- ❑ A possible heuristic function for this search space:  
#regressed goals that are not true in StartState:

$$h = | \text{RegressedGoals} - \text{StartState} |$$

# Question

---

- ☐ Is this heuristic function for the blocks world optimistic? That is, does it satisfy the condition of admissibility theorem for best-first search?
  
- ☐ If not, can it be modified to become optimistic?

# Planning as Constraint Satisfaction

---

CSP for each planning horizon  $k$  (vary  $k$  as needed)

## □ Variables

- Create a variable for each literal and time  $0, \dots, k$
- Create a variable for each action and time  $0, \dots, k-1$

## □ Constraints

- State constraints: literals at time  $t$
- Precondition constraints: actions and states at time  $t$
- Effect constraints: actions at time  $t$ , literals at times  $t$  and  $t + 1$
- Action constraints: actions at time  $t$  (mutual exclusion)
- Initial state constraints: literals at time 0
- Goal constraints: literals at time  $k$

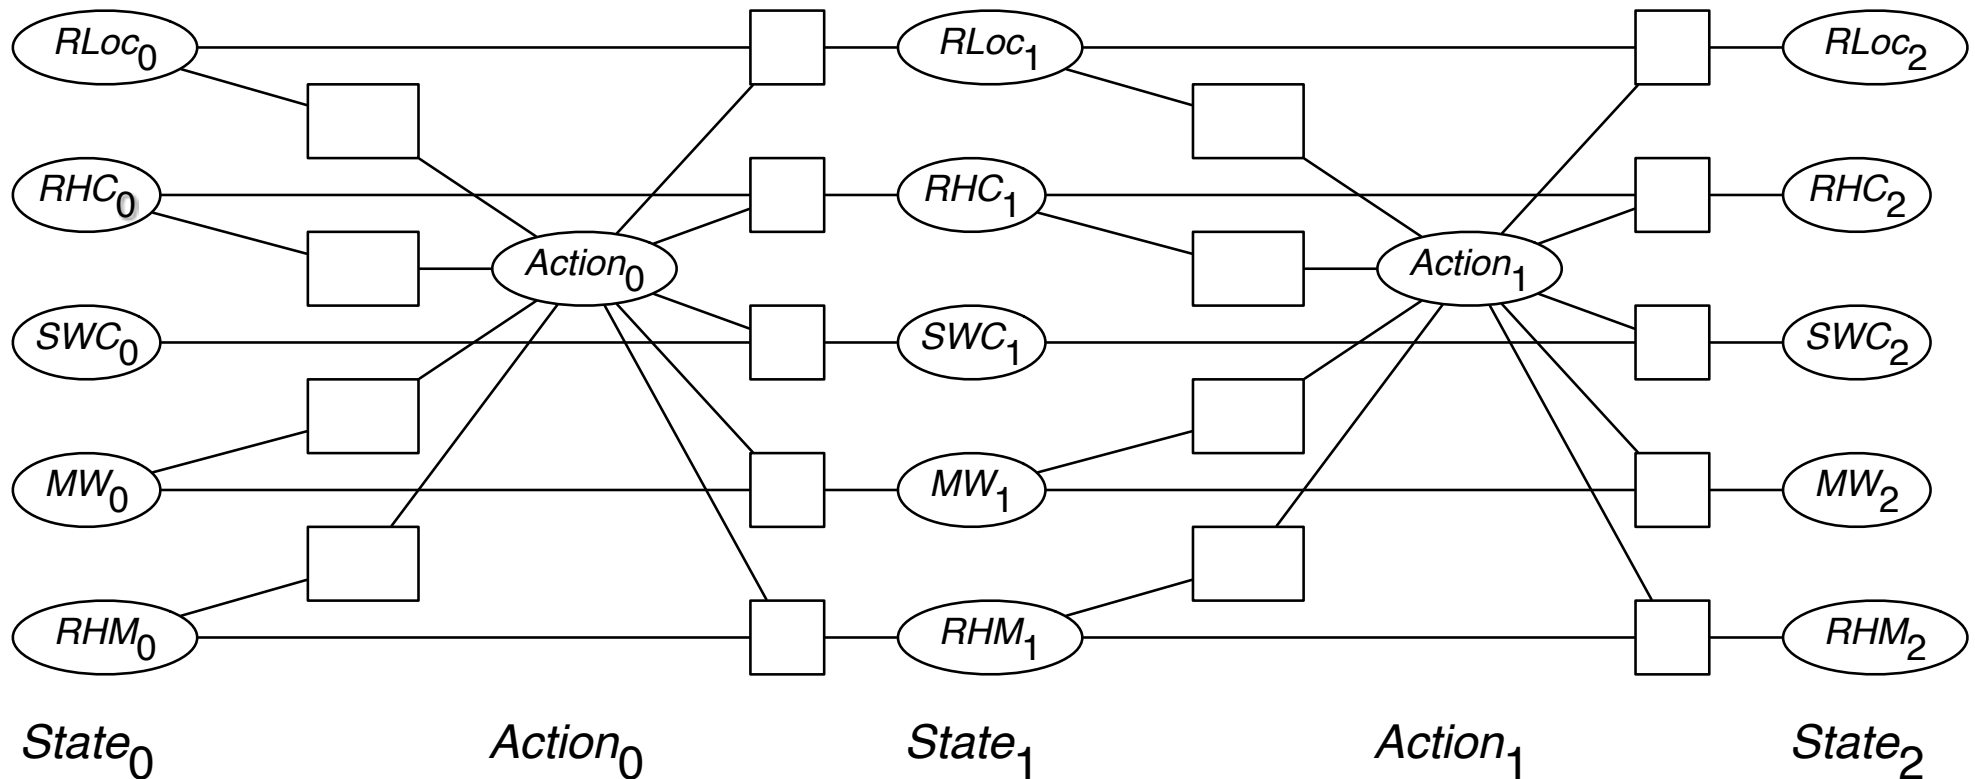


## The delivery robot CSP planner with factored actions

---

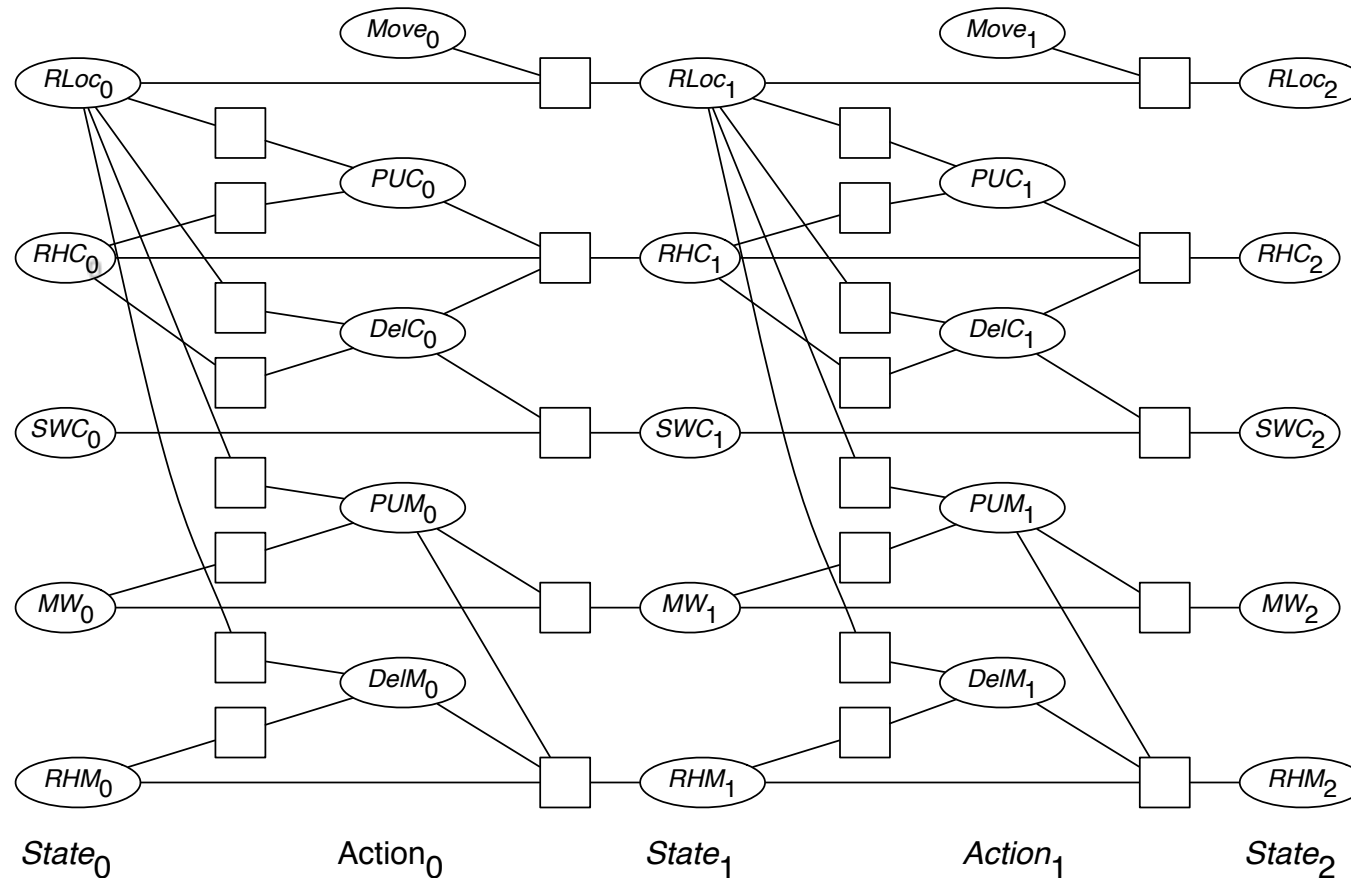
$RLoc_i$   
– Rob's location  
 $RHC_i$   
– Rob has coffee  
 $SWC_i$   
– Sam wants coffee  
 $MW_i$   
– Mail is waiting  
 $RHM_i$   
– Rob has mail  
 $Move_i$   
– Rob's move action  
 $PUC_i$   
– Rob picks up coffee  
 $DelC$   
– Rob delivers coffee  
 $PUM_i$   
– Rob picks up mail  
 $DelM_i$   
– Rob delivers mail

# Planning as Constraint Satisfaction



*CSP for planning horizon of 2.*

# Planning as Constraint Satisfaction



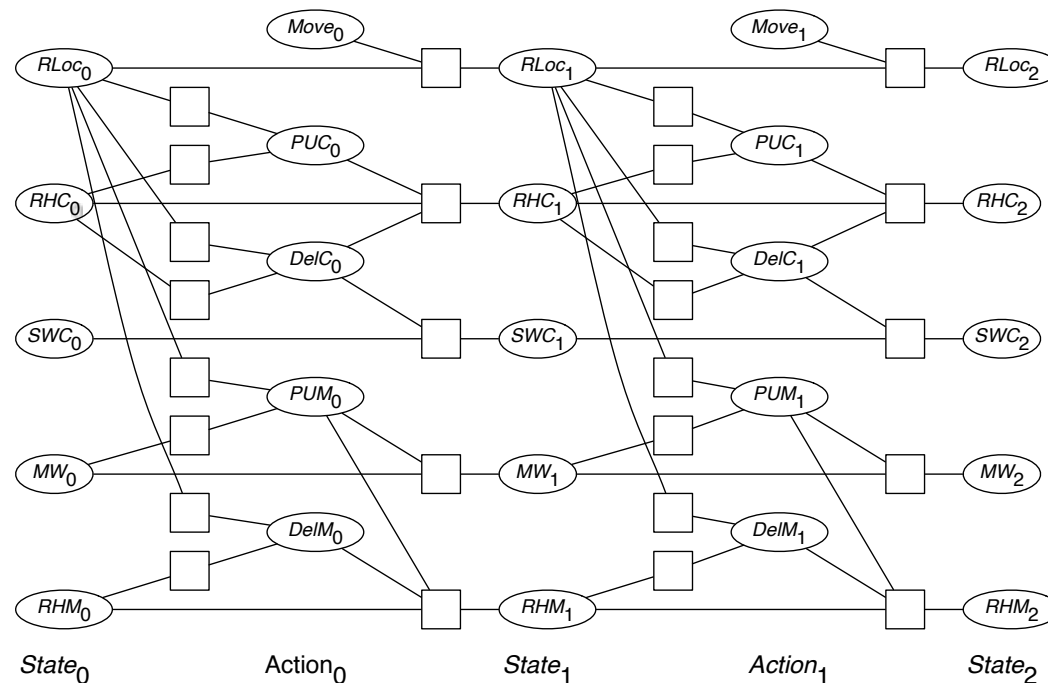
*CSP for planning horizon of 2.*

There can be an extra set of constraints **mutex constraints** - to specify which action features cannot co-occur.

# Planning as Constraint Satisfaction

The agent can be seen as doing more than one action in a single stage.

- For some of the actions at the same stage, the robot can do them in any order, such as delivering coffee and delivering mail.
- Some of the actions at the same stage need to be carried out in a particular order, for example, the agent must move after the other actions.



# Summary

---

- ❑ Planning is the process of choosing a sequence of actions to achieve a goal.
- ❑ An action is a partial function from a state to a state.
- ❑ Two representations for actions that exploit structure in states are
  - the STRIPS representation, which is an action-centric representation,
  - the feature-based representation of actions, which is a feature-centric representation.
  - The feature-based representation is more powerful than the STRIPS representation; it can represent anything representable in STRIPS, but can also represent conditional effects.

# Summary

---

- ❑ A forward planner searches in the state space from the initial state to a goal state.
- ❑ A regression planner searches backwards from the goal, where each node in the search space is a subgoal to be achieved.

# Summary

---

- ❑ Reasoning about action interesting from philosophical point of view
- ❑ Recent advances in planning give great improvements in efficiency
- ❑ Planning makes use of CSP framework with heuristics
- ❑ Multi-agent systems, dynamic worlds much more complex

# References

---

- ❑ Poole & Mackworth, Artificial Intelligence: Foundations of Computational Agents, Chapter 6
- ❑ I. Bratko, *Prolog Programming for Artificial Intelligence*, 4th Edition, Chapter 17.